

Heart Stroke

October 30, 2023

```
[362]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from numpy import mean
from numpy import std
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import MinMaxScaler
from sklearn.inspection import permutation_importance
from statsmodels.graphics.gofplots import qqplot
from scipy.stats import shapiro
```

```
[328]: df = pd.read_csv(f'D:\machine learning(sharif)\my project/
↳healthcare-dataset-stroke-data.csv')
df.head()
```

```
[328]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	

```

stroke
0      1
1      1
2      1
3      1
4      1

```

```
[329]: df.drop('id', axis=1, inplace=True)
```

```
[330]: df.head(5)
```

```
[330]:   gender  age  hypertension  heart_disease  ever_married  work_type \
0   Male  67.0             0             1         Yes      Private
1  Female  61.0             0             0         Yes  Self-employed
2   Male  80.0             0             1         Yes      Private
3  Female  49.0             0             0         Yes      Private
4  Female  79.0             1             0         Yes  Self-employed
```

```

Residence_type  avg_glucose_level  bmi  smoking_status  stroke
0      Urban             228.69  36.6  formerly smoked      1
1      Rural             202.21   NaN  never smoked      1
2      Rural             105.92  32.5  never smoked      1
3      Urban             171.23  34.4      smokes      1
4      Rural             174.12  24.0  never smoked      1

```

0.1 Data Preprocessing

```
[331]: df.describe()
```

```
[331]:   count  age  hypertension  heart_disease  avg_glucose_level  \
count  5110.000000  5110.000000  5110.000000  5110.000000
mean    43.226614    0.097456    0.054012    106.147677
std     22.612647    0.296607    0.226063    45.283560
min      0.080000    0.000000    0.000000    55.120000
25%     25.000000    0.000000    0.000000    77.245000
50%     45.000000    0.000000    0.000000    91.885000
75%     61.000000    0.000000    0.000000   114.090000
max     82.000000    1.000000    1.000000   271.740000
```

```

count  bmi  stroke
count  4909.000000  5110.000000
mean    28.893237    0.048728
std      7.854067    0.215320
min     10.300000    0.000000
25%     23.500000    0.000000
50%     28.100000    0.000000
75%     33.100000    0.000000

```


```
max          97.600000      1.000000
```

```
[332]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 5110 non-null   object
1   age                   5110 non-null   float64
2   hypertension           5110 non-null   int64
3   heart_disease          5110 non-null   int64
4   ever_married           5110 non-null   object
5   work_type              5110 non-null   object
6   Residence_type         5110 non-null   object
7   avg_glucose_level      5110 non-null   float64
8   bmi                   4909 non-null   float64
9   smoking_status         5110 non-null   object
10  stroke                 5110 non-null   int64
dtypes: float64(3), int64(3), object(5)
memory usage: 439.3+ KB
```

```
[333]: sns.heatmap(df.isnull())
```

```
[333]: <AxesSubplot:>
```



```
[334]: df.isnull().sum()
```

```
[334]: gender          0
      age            0
      hypertension    0
      heart_disease   0
      ever_married    0
      work_type       0
      Residence_type  0
      avg_glucose_level 0
      bmi            201
      smoking_status  0
      stroke         0
      dtype: int64
```

```
[335]: df['bmi'].fillna(df['bmi'].mean(), inplace=True)
```

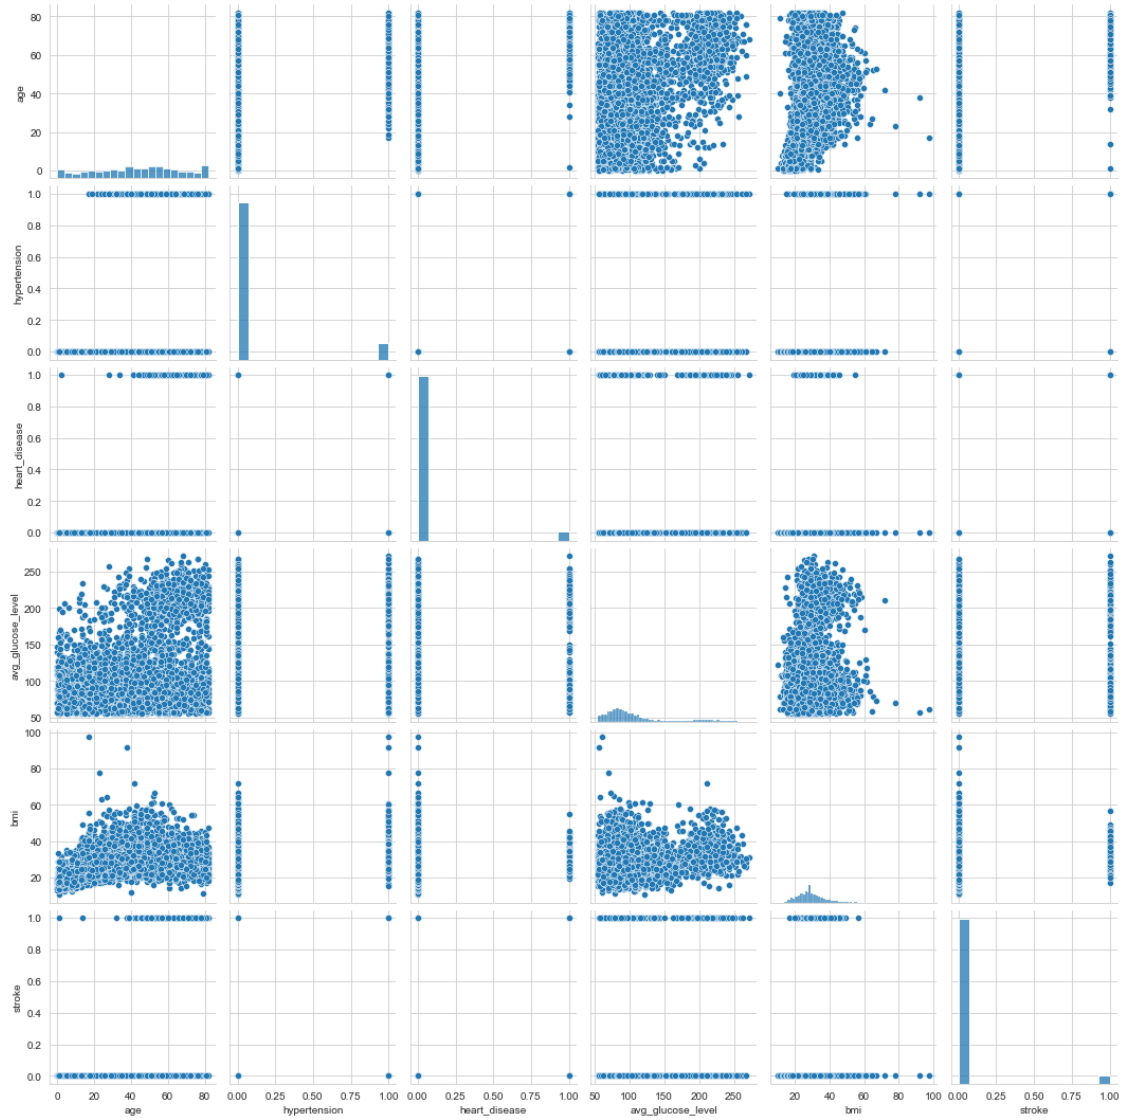
```
[336]: df['bmi'].isnull().sum()
```

```
[336]: 0
```


```
[337]: df.nunique()
```

```
[337]: gender          3  
age          104  
hypertension    2  
heart_disease   2  
ever_married    2  
work_type       5  
Residence_type  2  
avg_glucose_level 3979  
bmi            419  
smoking_status  4  
stroke          2  
dtype: int64
```

```
[338]: sns.set_style('whitegrid')  
sns.pairplot(df)  
plt.show()
```

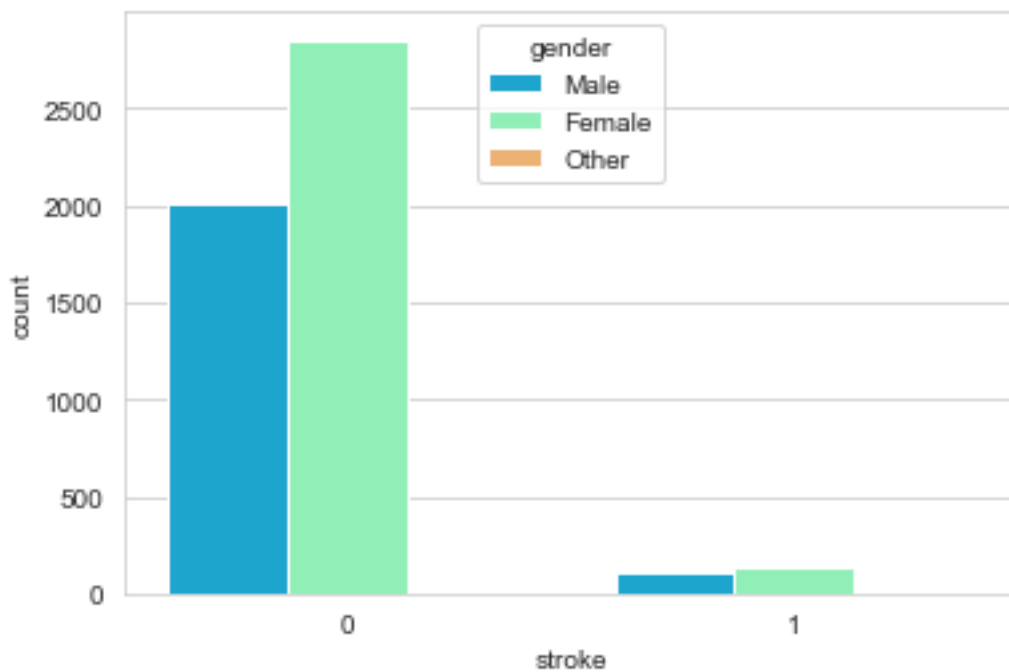


```
[339]: plt.figure(figsize = (30, 25))
sns.heatmap(df.corr(), annot = True, cmap="coolwarm")
plt.show()
```




```
[340]: sns.set_style('whitegrid')
sns.countplot(x='stroke',hue='gender',data=df,palette='rainbow')
```

```
[340]: <AxesSubplot:xlabel='stroke', ylabel='count'>
```




```
[341]: sns.set_style('whitegrid')
sns.countplot(x='stroke',hue='heart_disease',data=df,palette='rainbow')
```

```
[341]: <AxesSubplot:xlabel='stroke', ylabel='count'>
```



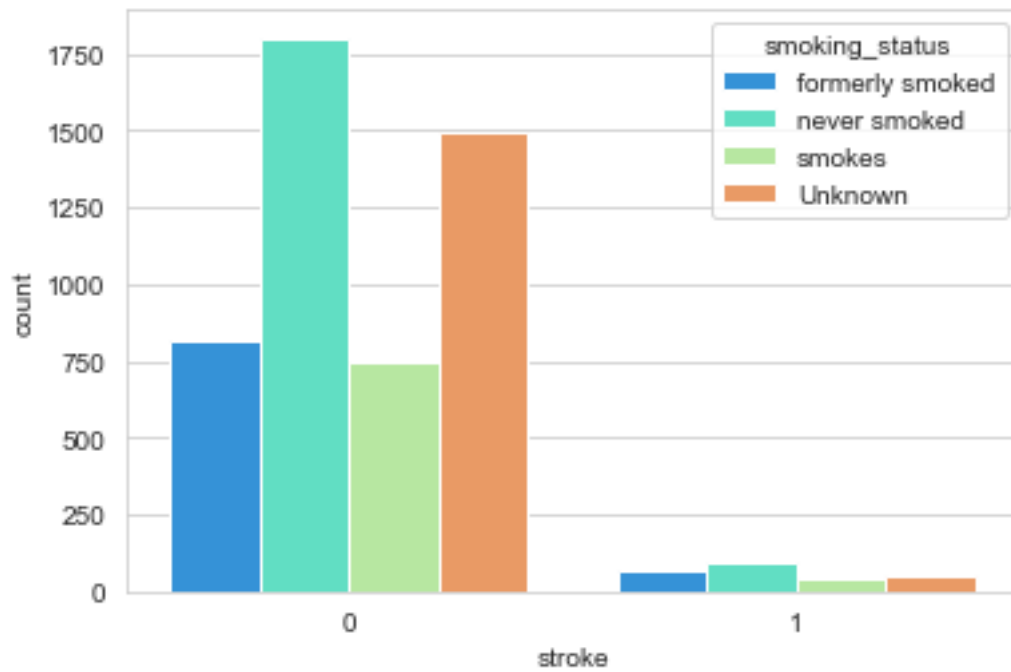

```
[342]: sns.set_style('whitegrid')
sns.countplot(x='stroke',hue='ever_married',data=df,palette='rainbow')
```

[342]: <AxesSubplot:xlabel='stroke', ylabel='count'>




```
[343]: sns.set_style('whitegrid')
sns.countplot(x='stroke',hue='smoking_status',data=df,palette='rainbow')
```

[343]: <AxesSubplot:xlabel='stroke', ylabel='count'>



```
[344]: sns.jointplot(y='avg_glucose_level',x='age',data=df,kind='scatter')
```


```
[344]: <seaborn.axisgrid.JointGrid at 0x226b71032e0>
```



0.2 Normality Test


```
[345]: df['age'].hist(bins=30,color='darkred',alpha=0.7)
```


```
[345]: <AxesSubplot:>
```



```
[346]: qqplot(df['age'],line='s')
```

[346]:





```
[347]: Statistics, p = shapiro(df['age'])
print(f'Statistics={Statistics:0.3f} p_value={p:0.3f}')

alpha = 0.05

if p >= alpha:
    print('Sample looks Gaussian (fail to reject H0)')
else:
    print('Sample does not look Gaussian (reject H0)')
```

Statistics=0.967 p_value=0.000

Sample does not look Gaussian (reject H0)


C:\Users\USER\anaconda3\lib\site-packages\scipy\stats\morestats.py:1760:

UserWarning: p-value may not be accurate for N > 5000.

warnings.warn("p-value may not be accurate for N > 5000.")


```
[348]: df['avg_glucose_level'].hist(bins=30,color='darkred',alpha=0.7)
```


```
[348]: <AxesSubplot:>
```



```
[349]: qqplot(df['avg_glucose_level'],line='s')
```

[349]:





```
[350]: Statistics, p = shapiro(df['avg_glucose_level'])
print(f'Statistics={Statistics:0.3f} p_value={p:0.3f}')

alpha = 0.05

if p >= alpha:
    print('Sample looks Gaussian (fail to reject H0)')
else:
    print('Sample does not look Gaussian (reject H0)')
```

Statistics=0.806 p_value=0.000

Sample does not look Gaussian (reject H0)


C:\Users\USER\anaconda3\lib\site-packages\scipy\stats\morestats.py:1760:

UserWarning: p-value may not be accurate for N > 5000.

warnings.warn("p-value may not be accurate for N > 5000.")


```
[351]: df['bmi'].hist(bins=30,color='darkred',alpha=0.7)
```


```
[351]: <AxesSubplot:>
```



```
[352]: qqplot(df['bmi'],line='s')
```

```
[352]:
```





0.3 Data Transform

```
[353]: df['ever_married'].replace({'Yes':1, 'No':0}, inplace=True)
df['gender'].replace({'Male':1, 'Female':0, 'Other':2}, inplace=True)
df['Residence_type'].replace({'Urban':1, 'Rural':0}, inplace=True)
df['smoking_status'].replace({'formerly smoked':0, 'never smoked':1, 'smokes':
    ↪2, 'Unknown':3}, inplace=True)
df['work_type'].replace({'Private':0, 'Self-employed':1, 'children':2,
    ↪'Govt_job':3, 'Never_worked':4}, inplace=True)
```

```
[354]: df.head()
```

```
[354]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	\
0	1	67.0	0	1	1	0	
1	0	61.0	0	0	1	1	
2	1	80.0	0	1	1	0	
3	0	49.0	0	0	1	0	
4	0	79.0	1	0	1	1	

	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	1	228.69	36.600000	0	1
1	0	202.21	28.893237	1	1
2	0	105.92	32.500000	1	1
3	1	171.23	34.400000	2	1
4	0	174.12	24.000000	1	1

```
[355]: X, y = df.drop('stroke', axis=1), df['stroke']
print(X.shape, y.shape)
numerical_ix = X.select_dtypes(include=['int64', 'float64', "float32"]).columns
print("numerical_ix: ",numerical_ix)
categorical_ix = X.select_dtypes(include=['object', 'bool']).columns
print("categorical_ix: ",categorical_ix)
t = [('cat', OneHotEncoder(), categorical_ix), ('num', MinMaxScaler(),
↪numerical_ix)]
col_transform = ColumnTransformer(transformers=t)
```

```
(5110, 10) (5110,)
numerical_ix: Index(['gender', 'age', 'hypertension', 'heart_disease',
'ever_married',
'work_type', 'Residence_type', 'avg_glucose_level', 'bmi',
'smoking_status'],
dtype='object')
categorical_ix: Index([], dtype='object')
```

0.4 Model Training and Testing

0.4.1 Logistic Regression

```
[356]: model=LogisticRegression()
cv = RepeatedStratifiedKFold(n_splits=10,n_repeats=5, random_state=1)
scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv,n_jobs=-1)
Accuracy_lg=mean(scores)
print('Accuracy_lg: %.3f (%.3f)' % (mean(scores), std(scores)))
```

```
Accuracy_lg: 0.951 (0.002)
```

0.4.2 SVM

```
[357]: model=SVC()
pipeline = Pipeline(steps=[('prep', col_transform), ('m', model)])
cv = RepeatedStratifiedKFold(n_splits=10,n_repeats=5, random_state=1)
scores = cross_val_score(pipeline, X, y, scoring='accuracy', cv=cv,n_jobs=-1)
Accuracy_svm=mean(scores)
print('Accuracy_svm: %.3f (%.3f)' % (mean(scores), std(scores)))
```

```
Accuracy_svm: 0.951 (0.001)
```

0.4.3 K-Nearest Neighbors (KNN)

```
[358]: model=KNeighborsClassifier()
pipeline = Pipeline(steps=[('prep', col_transform), ('m', model)])
cv = RepeatedStratifiedKFold(n_splits=10,n_repeats=5, random_state=1)
scores = cross_val_score(pipeline, X, y, scoring='accuracy', cv=cv,n_jobs=-1)
Accuracy_knn=mean(scores)
print('Accuracy_knn: %.3f (%.3f)' % (mean(scores), std(scores)))
```

Accuracy_knn: 0.949 (0.004)

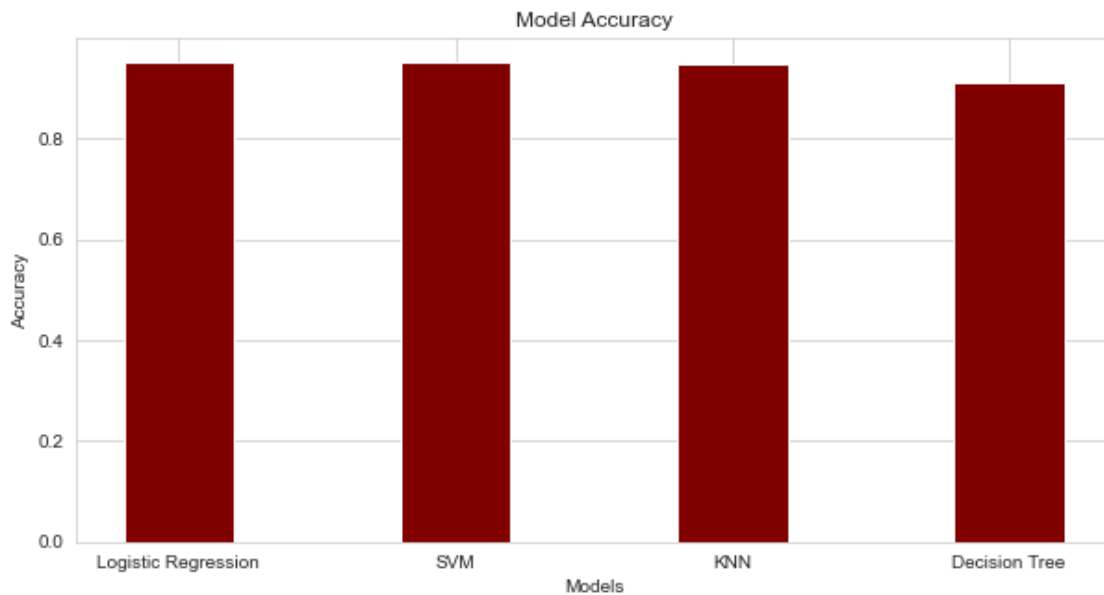
0.4.4 Decision Tree Classifier

```
[359]: model=DecisionTreeClassifier()
pipeline = Pipeline(steps=[('prep',col_transform), ('m', model)])
cv = RepeatedStratifiedKFold(n_splits=10,n_repeats=5, random_state=1)
scores = cross_val_score(pipeline, X, y, scoring='accuracy', cv=cv,n_jobs=-1)
Accuracy_dt=mean(scores)
print('Accuracy_dt: %.3f (%.3f)' % (mean(scores), std(scores)))
```

Accuracy__dt: 0.911 (0.010)

0.5 Model Comparison


```
[360]: models = ['Logistic Regression', 'SVM', 'KNN','Decision Tree']
accuracy = [Accuracy_lg, Accuracy_svm, Accuracy_knn,Accuracy_dt]
plt.figure(figsize=(10,5))
plt.bar(models, accuracy, color = 'Maroon', width = 0.4)
plt.xlabel('Models')
plt.ylabel('Accuracy')
plt.title('Model Accuracy')
plt.show()
```



0.5.1 Permutation Feature Importance

```
[361]: model = KNeighborsClassifier()
model.fit(X, y)
results = permutation_importance(model, X, y, scoring='accuracy')
importance = results.importances_mean
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))
plt.bar([x for x in range(len(importance))], importance)
plt.show()
```

```
Feature: 0, Score: 0.00043
Feature: 1, Score: 0.00595
Feature: 2, Score: -0.00023
Feature: 3, Score: -0.00027
Feature: 4, Score: -0.00004
Feature: 5, Score: -0.00043
Feature: 6, Score: -0.00035
Feature: 7, Score: 0.00583
Feature: 8, Score: 0.00290
Feature: 9, Score: 0.00047
```



0.6 Conclusion

The model accuracies of Logistic Regression and SVM are quite similar 95.1 %. The accuracy of KNN and Decision Tree Classifier are 94.9 % and 91 % So, we can use any of these models to predict the heart stroke.

The relationship of different features was depicted, but at the end, based on the KNN model, the characteristics that had the greatest impact on the prediction were identified. These features are respectively: 1-age 2-avg_glucose_level 3-bmi

[]: